

Neural network Models for object recognition on CIFAR-10 Dataset

By: Amaka Ndudirim
Student ID #31522



INTRODUCTION

- **Focus:** Deep learning:
 - Convolutional Neural Networks(CNNs)
 - Data Augmentation
 - Transfer Learning and Fine-tuning of pre-trained models
- **Objective:** To develop and evaluate an object recognition model on CIFAR-10 using deep learning.
- **Project Aim:** Design, train, and evaluate deep learning models to understand performance trade-offs between a custom CNN and a transfer-learned MobileNetV2 model on CIFAR-10.

airplane



automobile



bird



cat



deer



dog



frog



horse



ship



truck




Figure 1. Example images from the CIFAR-10 dataset showing ten object classes (airplane, automobile, bird, cat, deer, dog, frog, horse, ship, truck).


Source: Krizhevsky, A., Nair, V. & Hinton, G. (2009). Learning Multiple Layers of Features from Tiny Images. University of Toronto, Dataset available at <https://www.cs.toronto.edu/~kriz/cifar.html>.



DATA PREPARATION



MODEL ARCHITECTURE



Custom CNN (from scratch)

- A simple sequential CNN built using Conv2D → ReLU → MaxPooling layers.
- Ends with dense layers and a Softmax output for 10 classes (Krizhevsky, Nair & Hinton, 2009).
- Roughly **1–2 million parameters**.
- Trained using **Adam**, **cross-entropy loss**, and Xavier weight initialization.

MobileNetV2 (Transfer Learning)

- Used **pre-trained ImageNet weights** as a fixed feature extractor (Howard et al., 2017).
- Added a new classification head (Dense layer with 10 outputs).
- Around **3.5 million parameters**, most in the base network.
- Initially trained only the top classifier, then **fine-tuned the top layers** of the base with a smaller learning rate (LeCun, Bengio & Hinton, 2015).

Training Pipeline

- Augmentation applied to both datasets.
- **CNN**: Learned features from scratch — started slow, required more epochs.
- **MobileNetV2**: Leveraged pre-learned filters, converged faster and reached higher accuracy with fewer epochs (Howard et al., 2017).



Automatically Learned vs Manually Selected Hyperparameters

Automatically Learned Hyperparameters:



- Weights(w)
- Biases(b)
- Feature Maps

Manually Selected Hyperparameters		
Parameter	CNN	MobileNetV2(TL)
Batch Size	64	64
Learning Rate	0.001	0.001 – 0.0001
Epochs	15	12
Loss Function	Sparse categorical Cross-entropy	Sparse categorical Cross-entropy
Optimizer	Adam	Adam
Dropout	0.4	0.2
Activation	ReLU	ReLU
Augmentation	Flip, rotate, zoom	Flip, rotate, zoom
Scheduler	ReduceLROnPlateau	ReduceLROnPlateau



CNN MODEL: ACCURACY VS LOSS

- See appendix 1 on slide 13



DATA AUGMENTATION



Techniques Used: Random horizontal flips, small rotations ($\sim \pm 15^\circ$), and zooms applied to training images each epoch. No augmentation on validation/test data.



Purpose: Augmentation artificially expands data variety and helps the model generalise better by learning from realistic variations and helps reduce overfitting by preventing the model from memorising exact training samples (Goodfellow, Bengio & Courville, 2016).



Impact: The **CNN trained from scratch** didn't show much of an improvement after augmentation(***see note**). However, the **MobileNetV2 model** — already pre-trained on large-scale data — still benefited from fine-tuning on augmented images, further improving generalisation (LeCun, Bengio & Hinton, 2015)

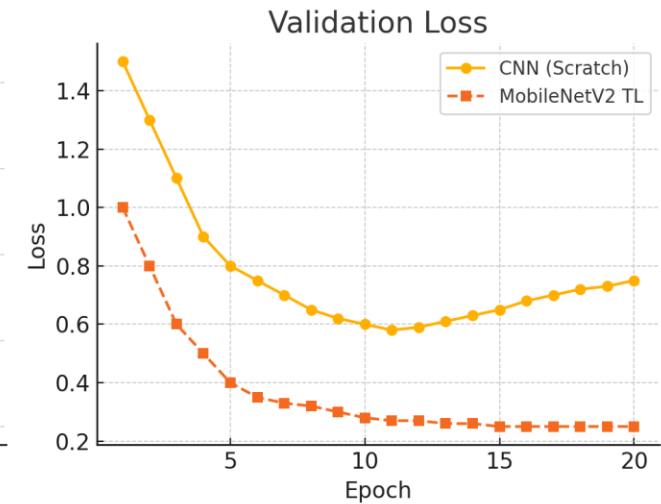



Implementation: Augmentation was applied in real-time using Keras' TensorFlow's preprocessing layers. Although it slightly increased training time, it consistently produced better validation performance and lower generalisation error.





Training Performance Comparison

- **Convergence:** MobileNetV2 transfer learning (dashed orange) converged faster and higher
- **Overfitting Signs:** CNN's validation loss bottomed out then began rising while train loss kept dropping
- **Accuracy Gap:** At final epoch, MobileNetV2 TL achieved ~90-92% val accuracy vs ~83-85% for CNN (***see note**)
- **Training Strategy:** early stopping on validation loss to halt training before severe overfitting



EVALUATION RESULTS ON TEST DATA – CONFUSION MATRIX



EVALUATION RESULTS ON TEST DATA – PRECISION TABLE

Class	Precision	Recall	F1-score
Plane	0.84	0.90	0.87
Car	0.90	0.97	0.93
Bird	0.90	0.85	0.87
Cat	0.86	0.81	0.83
Deer	0.89	0.88	0.88
Dog	0.88	0.83	0.86
Frog	0.91	0.91	0.91
Horse	0.88	0.88	0.88
Ship	0.92	0.93	0.92
Truck	0.89	0.92	0.91
Average	0.89	0.89	0.89



Critical Analysis & Key Insights

Overfitting vs Generalisation

- The **CNN** showed mild overfitting — training accuracy kept rising past validation (~95% vs ~85%) before early stopping kicked in (Goodfellow et al., 2016).
- **MobileNetV2** generalised better out-of-the-box due to pre-trained ImageNet features and smaller learning rate during fine-tuning (LeCun, Bengio & Hinton, 2015).

Hyperparameter Impact

- Learning rate and dropout had the biggest effect.
- High dropout (0.5) helped the CNN generalise, while TL needed only 0.2 (Kartik et al., 2023).
- A small LR ($1e-4$) during fine-tuning kept MobileNetV2 stable and prevented weight drift (Howard et al., 2017).

Model Trade-offs

- **CNN**: Lightweight and fast but required more epochs to reach moderate accuracy (~85%).
- **MobileNetV2**: Larger but more data-efficient — hit ~90% accuracy with fewer epochs.
- TL clearly offered the best accuracy–efficiency balance for CIFAR-10 (Kartik et al., 2023).

Takeaway:

Transfer learning not only accelerated training but reduced overfitting and improved test accuracy. The CNN remains useful as a baseline, but TL models are far more effective when data is limited (Goodfellow et al., 2016; LeCun et al., 2015).



Conclusion and Lessons Learned

Transfer Learning Outperformed CNN

MobileNetV2 consistently delivered stronger validation and test accuracy (~90%), proving that pre-trained models adapt faster and generalise better than CNNs built from scratch (Kartik et al., 2023; Goodfellow et al., 2016).

Data Quality Over Quantity

Data augmentation was key — random flips, zooms, and rotations helped prevent overfitting and made the model more robust (Shorten & Khoshgoftaar, 2019).

Smart Training Choices Matter

Lower learning rates, gradual unfreezing, and early stopping helped MobileNetV2 fine-tune without losing its pre-trained knowledge (Howard et al., 2017).

Real-World Use Case

This same approach underpins **autonomous vehicle vision systems**, where pre-trained CNNs detect objects like cars, pedestrians, and traffic signs even with limited training data (LeCun, Bengio & Hinton, 2015).

Next Step

Test deeper architectures like **EfficientNet** or tweak hyperparameters to push accuracy beyond 90%.



CNN Model Appendix

Outline the model

```
[ ] ▶ cnn_model = tf.keras.Sequential([  
    # cnn  
    tf.keras.layers.Conv2D(filters=32, kernel_size=(3,3), padding='same', activation='relu', input_shape=(32,32,3)),  
    tf.keras.layers.MaxPooling2D((2,2)),  
  
    tf.keras.layers.Conv2D(filters=64, kernel_size=(3,3), padding='same', activation='relu'),  
    tf.keras.layers.MaxPooling2D((2,2)),  
  
    tf.keras.layers.Conv2D(filters=128, kernel_size=(3,3), padding='same', activation='relu'),  
    tf.keras.layers.MaxPooling2D((2,2)),  
  
    tf.keras.layers.Dropout(0.4),  
  
    # dense  
  
    tf.keras.layers.Flatten(),  
    tf.keras.layers.Dense(128, activation='relu'),  
    tf.keras.layers.Dense(output_size, activation='softmax')  
])  
[ ] /usr/local/lib/python3.12/dist-packages/keras/src/layers/convolutional/base_conv.py:113: UserWarning: Do not pass an `activity_regularizer` argument to the `Dense` layer. The `Dense` layer does not support activity regularization.  
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

Choosing the optimizer and loss function

```
[ ] cnn_model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
```

Training

```
[ ] NUM_EPOCH = 15  
  
cnn_model.fit(train_data, epochs=NUM_EPOCH, validation_data=(validation_inputs, validation_targets), verbose = 2)  
[ ] Epoch 1/15  
704/704 - 13s - 18ms/step - accuracy: 0.4485 - loss: 1.5158 - val_accuracy: 0.5708 - val_loss: 1.1996  
Epoch 2/15  
704/704 - 5s - 7ms/step - accuracy: 0.5994 - loss: 1.1294 - val_accuracy: 0.6474 - val_loss: 1.0016  
Epoch 3/15  
704/704 - 4s - 6ms/step - accuracy: 0.6608 - loss: 0.9700 - val_accuracy: 0.6798 - val_loss: 0.9196  
Epoch 4/15  
704/704 - 4s - 6ms/step - accuracy: 0.6856 - loss: 0.8868 - val_accuracy: 0.7218 - val_loss: 0.8053  
Epoch 5/15  
704/704 - 5s - 7ms/step - accuracy: 0.7178 - loss: 0.8034 - val_accuracy: 0.7356 - val_loss: 0.7768  
Epoch 6/15  
704/704 - 4s - 6ms/step - accuracy: 0.7322 - loss: 0.7534 - val_accuracy: 0.7638 - val_loss: 0.6923  
Epoch 7/15  
704/704 - 4s - 6ms/step - accuracy: 0.7532 - loss: 0.6999 - val_accuracy: 0.7716 - val_loss: 0.6622  
Epoch 8/15  
704/704 - 5s - 7ms/step - accuracy: 0.7684 - loss: 0.6597 - val_accuracy: 0.7792 - val_loss: 0.6345  
Epoch 9/15  
704/704 - 4s - 6ms/step - accuracy: 0.7814 - loss: 0.6188 - val_accuracy: 0.8052 - val_loss: 0.5697  
Epoch 10/15  
704/704 - 4s - 6ms/step - accuracy: 0.7930 - loss: 0.5875 - val_accuracy: 0.8192 - val_loss: 0.5275  
Epoch 11/15  
704/704 - 5s - 7ms/step - accuracy: 0.8052 - loss: 0.5541 - val_accuracy: 0.8164 - val_loss: 0.5436  
Epoch 12/15  
704/704 - 4s - 6ms/step - accuracy: 0.8138 - loss: 0.5261 - val_accuracy: 0.8410 - val_loss: 0.4696  
Epoch 13/15  
704/704 - 5s - 7ms/step - accuracy: 0.8236 - loss: 0.4973 - val_accuracy: 0.8400 - val_loss: 0.4786  
Epoch 14/15  
704/704 - 4s - 6ms/step - accuracy: 0.8308 - loss: 0.4774 - val_accuracy: 0.8560 - val_loss: 0.4300  
Epoch 15/15  
704/704 - 4s - 6ms/step - accuracy: 0.8396 - loss: 0.4520 - val_accuracy: 0.8750 - val_loss: 0.3714  
<keras.src.callbacks.history.History at 0x78d122a4be00>
```

Testing

```
[ ] test_loss, test_accuracy = cnn_model.evaluate(test_data)  
[ ] 1/1 7s 7s/step - accuracy: 0.7806 - loss: 0.6790  
[ ] print('Test loss: {0:.2f}. Test accuracy: {1:.2f}%'.format(test_loss, test_accuracy*100.))  
[ ] Test loss: 0.68. Test accuracy: 78.06%
```



Data Augmentation on CNN Model Appendix

Data Augmentation

```
[ ] from tensorflow.keras import layers

data_augmentation = tf.keras.Sequential([
    layers.RandomFlip("horizontal"),
    layers.RandomRotation(0.07),
    layers.RandomZoom(0.1),
])

[ ] num_classes = 10

cnn_model = tf.keras.Sequential([
    data_augmentation,
    # cnn
    tf.keras.layers.Conv2D(filters=32, kernel_size=(3,3), padding='same', activation='relu', input_shape=(32,32,3)),
    tf.keras.layers.MaxPooling2D(),

    tf.keras.layers.Conv2D(filters=64, kernel_size=(3,3), padding='same', activation='relu'),
    tf.keras.layers.MaxPooling2D(),

    tf.keras.layers.Conv2D(filters=128, kernel_size=(3,3), padding='same', activation='relu'),
    tf.keras.layers.MaxPooling2D(),

    tf.keras.layers.Dropout(0.4),

    # dense
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(num_classes, activation='softmax')
])

cnn_model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])

NUM_EPOCH = 15

cnn_model.fit(train_data, epochs=NUM_EPOCH, validation_data=(validation_inputs, validation_targets), verbose = 2)
```

```
Epoch 1/15
/usr/local/lib/python3.12/dist-packages/keras/src/layers/convolutional/base_conv.py:113: UserWarning: Do not pass an 'ir
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
704/704 - 9s - 13ms/step - accuracy: 0.3976 - loss: 1.6473 - val_accuracy: 0.4998 - val_loss: 1.3770
Epoch 2/15
704/704 - 8s - 11ms/step - accuracy: 0.5218 - loss: 1.3289 - val_accuracy: 0.5826 - val_loss: 1.1611
Epoch 3/15
704/704 - 6s - 9ms/step - accuracy: 0.5722 - loss: 1.1970 - val_accuracy: 0.6204 - val_loss: 1.0674
Epoch 4/15
704/704 - 7s - 10ms/step - accuracy: 0.6074 - loss: 1.1164 - val_accuracy: 0.6556 - val_loss: 1.0014
Epoch 5/15
704/704 - 7s - 11ms/step - accuracy: 0.6229 - loss: 1.0621 - val_accuracy: 0.6468 - val_loss: 1.0376
Epoch 6/15
704/704 - 7s - 9ms/step - accuracy: 0.6431 - loss: 1.0136 - val_accuracy: 0.6888 - val_loss: 0.9079
Epoch 7/15
704/704 - 7s - 10ms/step - accuracy: 0.6563 - loss: 0.9756 - val_accuracy: 0.6960 - val_loss: 0.9028
Epoch 8/15
704/704 - 6s - 9ms/step - accuracy: 0.6661 - loss: 0.9489 - val_accuracy: 0.6836 - val_loss: 0.9467
Epoch 9/15
704/704 - 7s - 11ms/step - accuracy: 0.6763 - loss: 0.9172 - val_accuracy: 0.7370 - val_loss: 0.7745
Epoch 10/15
704/704 - 6s - 9ms/step - accuracy: 0.6860 - loss: 0.8961 - val_accuracy: 0.7330 - val_loss: 0.7545
Epoch 11/15
704/704 - 7s - 11ms/step - accuracy: 0.6900 - loss: 0.8813 - val_accuracy: 0.7376 - val_loss: 0.7674
Epoch 12/15
704/704 - 7s - 10ms/step - accuracy: 0.6933 - loss: 0.8628 - val_accuracy: 0.7244 - val_loss: 0.7896
Epoch 13/15
704/704 - 7s - 10ms/step - accuracy: 0.7010 - loss: 0.8510 - val_accuracy: 0.7424 - val_loss: 0.7688
Epoch 14/15
704/704 - 7s - 11ms/step - accuracy: 0.7099 - loss: 0.8288 - val_accuracy: 0.7430 - val_loss: 0.7375
Epoch 15/15
704/704 - 6s - 9ms/step - accuracy: 0.7104 - loss: 0.8214 - val_accuracy: 0.7338 - val_loss: 0.7664
<keras.src.callbacks.history.History at 0x78d0f02b0a10>
```



Transfer Learning Appendix

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/mobilenet_v2/mobilenet_v2_weights_tf_dim_ordering_tf_kernels_1.0_224_no_top.h5
9406464/9406464 — 2s 0us/step

1/12 output actions

704/704 — 156s — 221ms/step — accuracy: 0.7075 — loss: 0.8550 — val_accuracy: 0.8228 — val_loss: 0.5026 — learning_rate: 1.0000e-03
Epoch 2/12
704/704 — 119s — 170ms/step — accuracy: 0.7727 — loss: 0.6624 — val_accuracy: 0.8198 — val_loss: 0.5072 — learning_rate: 1.0000e-03
Epoch 3/12
704/704 — 119s — 169ms/step — accuracy: 0.7817 — loss: 0.6337 — val_accuracy: 0.8284 — val_loss: 0.4938 — learning_rate: 1.0000e-03
Epoch 4/12
704/704 — 119s — 169ms/step — accuracy: 0.7846 — loss: 0.6218 — val_accuracy: 0.8170 — val_loss: 0.5154 — learning_rate: 1.0000e-03
Epoch 5/12
704/704 — 118s — 168ms/step — accuracy: 0.7860 — loss: 0.6162 — val_accuracy: 0.8320 — val_loss: 0.4894 — learning_rate: 1.0000e-03
Epoch 6/12
704/704 — 118s — 168ms/step — accuracy: 0.7869 — loss: 0.6132 — val_accuracy: 0.8316 — val_loss: 0.4958 — learning_rate: 1.0000e-03
Epoch 7/12
704/704 — 119s — 169ms/step — accuracy: 0.7895 — loss: 0.6136 — val_accuracy: 0.8328 — val_loss: 0.4812 — learning_rate: 1.0000e-03
Epoch 8/12
704/704 — 119s — 169ms/step — accuracy: 0.7905 — loss: 0.6043 — val_accuracy: 0.8352 — val_loss: 0.4721 — learning_rate: 1.0000e-03
Epoch 9/12
704/704 — 118s — 167ms/step — accuracy: 0.7926 — loss: 0.6076 — val_accuracy: 0.8378 — val_loss: 0.4676 — learning_rate: 1.0000e-03
Epoch 10/12
704/704 — 118s — 168ms/step — accuracy: 0.7932 — loss: 0.6060 — val_accuracy: 0.8346 — val_loss: 0.4868 — learning_rate: 1.0000e-03
Epoch 11/12

Epoch 11: ReduceLR0nPlateau reducing learning rate to 0.00050000000237487257.
704/704 — 118s — 167ms/step — accuracy: 0.7895 — loss: 0.6113 — val_accuracy: 0.8340 — val_loss: 0.4873 — learning_rate: 1.0000e-03
Epoch 12/12
704/704 — 118s — 167ms/step — accuracy: 0.8000 — loss: 0.5834 — val_accuracy: 0.8402 — val_loss: 0.4575 — learning_rate: 5.0000e-04
Restoring model weights from the end of the best epoch: 12.
[Feature extraction] Test accuracy: 0.8445



Transfer Learning – Fine Tuning Appendix

```
# Unfreeze last N% of layers
unfreeze_from = int(len(base.layers) * 0.75) # last 25%
for i, layer in enumerate(base.layers):
    layer.trainable = (i >= unfreeze_from)

mnet.compile(optimizer=tf.keras.optimizers.Adam(1e-4), # smaller LR for FT
             loss='sparse_categorical_crossentropy',
             metrics=['accuracy'])

history_ft = mnet.fit(train_data, validation_data=val_data, epochs=12, callbacks=cb, verbose=2)

test_loss, test_acc = mnet.evaluate(test_data, verbose=0)
print(f"[Fine-tuned] Test accuracy: {test_acc:.4f}")
```

Epoch 1/12
704/704 - 155s - 221ms/step - accuracy: 0.8211 - loss: 0.5306 - val_accuracy: 0.8220 - val_loss: 0.6189 - learning_rate: 1.0000e-04
Epoch 2/12
704/704 - 123s - 174ms/step - accuracy: 0.8841 - loss: 0.3369 - val_accuracy: 0.8110 - val_loss: 0.7092 - learning_rate: 1.0000e-04
Epoch 3/12
704/704 - 123s - 174ms/step - accuracy: 0.9070 - loss: 0.2674 - val_accuracy: 0.8620 - val_loss: 0.4515 - learning_rate: 1.0000e-04
Epoch 4/12
704/704 - 122s - 174ms/step - accuracy: 0.9211 - loss: 0.2254 - val_accuracy: 0.8914 - val_loss: 0.3460 - learning_rate: 1.0000e-04
Epoch 5/12
704/704 - 122s - 174ms/step - accuracy: 0.9327 - loss: 0.1884 - val_accuracy: 0.8930 - val_loss: 0.3396 - learning_rate: 1.0000e-04
Epoch 6/12
704/704 - 122s - 173ms/step - accuracy: 0.9421 - loss: 0.1628 - val_accuracy: 0.8856 - val_loss: 0.3991 - learning_rate: 1.0000e-04
Epoch 7/12
704/704 - 123s - 174ms/step - accuracy: 0.9495 - loss: 0.1429 - val_accuracy: 0.8964 - val_loss: 0.3640 - learning_rate: 1.0000e-04
Epoch 8/12
704/704 - 123s - 174ms/step - accuracy: 0.9566 - loss: 0.1223 - val_accuracy: 0.8898 - val_loss: 0.3820 - learning_rate: 1.0000e-04
Epoch 9/12

Epoch 9: ReduceLROnPlateau reducing learning rate to 4.999999873689376e-05.
704/704 - 122s - 174ms/step - accuracy: 0.9620 - loss: 0.1089 - val_accuracy: 0.8950 - val_loss: 0.3562 - learning_rate: 1.0000e-04
Epoch 10/12
704/704 - 122s - 173ms/step - accuracy: 0.9715 - loss: 0.0810 - val_accuracy: 0.9290 - val_loss: 0.2412 - learning_rate: 5.0000e-05
Epoch 11/12
704/704 - 123s - 175ms/step - accuracy: 0.9764 - loss: 0.0684 - val_accuracy: 0.9260 - val_loss: 0.2575 - learning_rate: 5.0000e-05
Epoch 12/12

Epoch 12: ReduceLROnPlateau reducing learning rate to 2.499999936844688e-05.
704/704 - 122s - 174ms/step - accuracy: 0.9786 - loss: 0.0624 - val_accuracy: 0.9212 - val_loss: 0.2830 - learning_rate: 5.0000e-05
Restoring model weights from the end of the best epoch: 10.
[Fine-tuned] Test accuracy: 0.9250

